
HW1 Report: A Three-Layer MLP from Scratch for EuroSAT RGB Land Cover Classification

Czh
123456@qq.com

Abstract

This report presents Homework 1 on building a three-layer multilayer perceptron (MLP) from scratch for EuroSAT RGB land cover classification. The entire training pipeline was implemented with NumPy, including forward propagation, manual backpropagation, softmax cross-entropy, stochastic gradient descent (SGD), learning-rate decay, weight decay, checkpoint selection, evaluation, staged hyperparameter search, and visualization. No automatic differentiation framework such as PyTorch, TensorFlow, or JAX was used. In the final experiment, a two-stage search strategy was applied: a 54-trial coarse grid search first identified a promising region of the hyperparameter space, and a 27-trial fine grid search then refined the configuration around the coarse optimum. The coarse stage systematically compared ReLU, Tanh, and Sigmoid, and showed a clear advantage for ReLU. The final model, trained with the best fine-search configuration, reached a best validation accuracy of 0.6788 at epoch 30 and a test accuracy of 0.6780. The report further analyzes the search strategy, result tables, confusion matrix, first-layer weights, and representative failure cases.

1 Introduction

The goal of this homework is to implement an image classifier for the EuroSAT RGB dataset using a three-layer MLP trained entirely from scratch. The task is multi-class land cover classification with 10 categories: AnnualCrop, Forest, HerbaceousVegetation, Highway, Industrial, Pasture, PermanentCrop, Residential, River, and SeaLake.

The main technical constraint of the assignment is that the learning procedure must not rely on automatic differentiation. Therefore, all trainable layers, activation functions, gradients, and optimization steps were implemented manually in NumPy. The completed project includes modular components for data loading, preprocessing, model definition, training, hyperparameter search, testing, and visualization under `code/`, while all generated artifacts are saved under `results/`.

2 Dataset and Preprocessing

The dataset used in this homework is EuroSAT_RGB/, organized as one subdirectory per class. Inspection of the local dataset shows that the raw RGB images are consistently 64×64 . The implemented loader reads these images, resizes them to a fixed spatial resolution for training, converts them to floating-point arrays, and normalizes pixel values to $[0, 1]$.

In the main experiment, the raw 64×64 images were resized to 32×32 , resulting in an input dimensionality of $32 \times 32 \times 3 = 3072$ after flattening. A fixed random seed of 42 was used throughout the pipeline for reproducibility. The dataset was split in a stratified manner into training, validation, and test subsets with ratios 70%, 15%, and 15%, respectively. The resulting split sizes are shown in Table 1.

Table 1: Dataset split and input configuration used in the final training pipeline. The raw images are 64×64 , while the MLP input is resized to 32×32 .

Item	Value
Number of classes	10
Raw image size	$64 \times 64 \times 3$
Training input size	$32 \times 32 \times 3$
Flattened input dimension	3072
Training samples	18,900
Validation samples	4,050
Test samples	4,050
Split ratio	70 / 15 / 15
Random seed	42

Table 2: Class distribution in the EuroSAT RGB dataset, read from the local EuroSAT_RGB/ directory.

Class	Number of Images
AnnualCrop	3,000
Forest	3,000
HerbaceousVegetation	3,000
Highway	2,500
Industrial	2,500
Pasture	2,000
PermanentCrop	2,500
Residential	3,000
River	2,500
SeaLake	3,000
Total	27,000

Channel-wise standardization was computed from the training split only and then applied to all subsets. Mini-batch iteration was implemented in NumPy with shuffling enabled for training.

The resize step is a practical design choice rather than a claim that the original data are low resolution. For a fully connected MLP, using the raw 64×64 images would increase the flattened input dimensionality from 3072 to 12288 and would enlarge the first-layer parameter count by a factor of four. This would substantially increase training cost, make hyperparameter search less feasible, and raise the risk of overfitting for a dense architecture that does not explicitly exploit local spatial structure. Resizing to 32×32 inevitably removes some fine texture information, but it provides a better trade-off between computational cost, numerical stability, and classification performance for a from-scratch NumPy MLP.

Table 2 summarizes the class distribution directly read from the dataset folders. The dataset is nearly balanced, although Highway, Industrial, PermanentCrop, River, and Pasture contain fewer images than AnnualCrop, Forest, Residential, and SeaLake. This mild imbalance is useful to mention in the report because it partially explains why some categories are slightly harder than others even before model-specific limitations are considered.

3 Model Architecture

The classifier is a three-layer MLP with the following structure:

$$\text{input} \rightarrow \text{hidden}_1 \rightarrow \text{activation} \rightarrow \text{hidden}_2 \rightarrow \text{activation} \rightarrow \text{output}.$$

For the final main training run, the network used hidden dimensions 768 and 384 with ReLU activation in both hidden layers. The output layer produced logits for the 10 land cover classes. The implementation supports ReLU, Tanh, and Sigmoid activations. In the updated systematic search, all three activations were considered in the coarse stage, after which the fine stage focused only on ReLU because it was the strongest activation family in the broader comparison.

Each linear layer stores its own parameters and intermediate cache for backpropagation. The backward pass explicitly computes gradients for weights and biases, and the optimizer updates parameters using SGD. This design follows the assignment requirement that the training process be implemented without any automatic differentiation engine.

4 Training Method

The training objective is softmax cross-entropy with L_2 regularization applied to the weight matrices. The optimization pipeline is shared by both the hyperparameter-search runs and the final full training run. In every experiment, the implementation uses mini-batch SGD, explicit manual backpropagation, learning-rate decay, and model selection by validation accuracy.

For each candidate configuration evaluated during search, the code records training loss, validation loss, validation accuracy, and learning rate at every epoch. The checkpoint with the best validation accuracy is saved automatically to `results/checkpoints/`. After the hyperparameter search determines the best configuration, the same training procedure is reused for the final 30-epoch run.

The final full-training configuration, selected after the two-stage search, is summarized below:

- batch size: 128
- epochs: 30
- hidden dimensions: 768 and 384
- activation: ReLU
- initial learning rate: 0.06
- learning-rate decay: 0.95 per epoch
- minimum learning rate: 10^{-4}
- weight decay: 10^{-5}

5 Hyperparameter Search Strategy

To make the hyperparameter study more systematic, the final experiment used a two-stage grid-search strategy rather than a single short random search. The goal was to first identify a promising region of the search space and then refine the candidate configuration before running the full 30-epoch training used for the final model.

Stage 1: coarse search. The coarse search explored a relatively broad grid over four hyperparameters:

- learning rate: $\{0.03, 0.01, 0.003\}$
- hidden dimension of the first layer: $\{128, 256, 512\}$
- weight decay: $\{0, 10^{-4}\}$
- activation: `{relu, tanh, sigmoid}`

The second hidden dimension was set to half of the first hidden dimension. This coarse stage evaluated $3 \times 3 \times 2 \times 3 = 54$ configurations, each trained for 10 epochs. The best coarse-search configuration was ReLU, learning rate 0.03, hidden dimensions 512 and 256, and zero weight decay, with a best validation accuracy of 0.6600. Sigmoid was clearly weaker than both ReLU and Tanh in this broader comparison, so it was not retained for the fine search.

Stage 2: fine search. The fine search then focused on the neighborhood of the coarse optimum. Since the coarse search strongly favored ReLU and larger hidden widths, while Sigmoid performed substantially worse and Tanh remained consistently behind ReLU, the fine stage fixed the activation to ReLU and searched:

- learning rate: $\{0.015, 0.03, 0.06\}$

- hidden dimension of the first layer: $\{256, 512, 768\}$
- weight decay: $\{0, 10^{-5}, 10^{-4}\}$
- activation: `{relu}`

This produced $3 \times 3 \times 3 \times 1 = 27$ configurations, each trained for 15 epochs. The best fine-search configuration used learning rate 0.06, hidden dimensions 768 and 384, weight decay 10^{-5} , and ReLU activation, reaching a best validation accuracy of 0.6711.

This staged strategy is the main practical contribution of the final experiment: instead of treating hyperparameter search as a small sanity check, it was turned into a structured procedure that first covered the broader design space and then refined the most promising region.

Table 3: Summary of the two-stage hyperparameter search strategy and the best configuration found at each stage.

Stage	Trials	Epochs / Trial	Search Space	Best Configuration	Best Val Acc.
Coarse	54	10	LR $\in \{0.03, 0.01, 0.003\}$, hidden $\in \{128, 256, 512\}$, WD $\in \{0, 10^{-4}\}$, act $\in \{\text{ReLU}, \text{Tanh}, \text{Sigmoid}\}$	LR = 0.03, hidden = 512/256, WD = 0, ReLU	0.6600
Fine	27	15	LR $\in \{0.015, 0.03, 0.06\}$, hidden $\in \{256, 512, 768\}$, WD $\in \{0, 10^{-5}, 10^{-4}\}$, act = ReLU	LR = 0.06, hidden = 768/384, WD = 10^{-5} , ReLU	0.6711

The final model was then trained for 30 epochs using the best fine-search configuration. In other words, the search stage did not directly define the final result; it provided the strategy and the parameter choice for the full training run.

Table 4: Grouped statistics from the coarse and fine searches. Values are mean best validation accuracy over all configurations sharing the same hyperparameter value.

Hyperparameter	Coarse Search	Fine Search
Learning rate	0.03 : 0.6245, 0.01 : 0.5961, 0.003 : 0.5525	0.03 : 0.6619, 0.06 : 0.6602, 0.015 : 0.6539
Hidden dimension	512 : 0.6040, 256 : 0.5924, 128 : 0.5767	768 : 0.6659, 512 : 0.6637, 256 : 0.6463
Weight decay	0 : 0.5913, 10^{-4} : 0.5908	0 : 0.6590, 10^{-5} : 0.6588, 10^{-4} : 0.6580
Activation	ReLU: 0.6019, Tanh: 0.5802, Sigmoid: 0.3915	ReLU only

Table 5: Top configurations from the two-stage search. The best results consistently favor ReLU, larger hidden widths, and relatively aggressive learning rates.

Stage	Trial	LR	Hidden ₁	Hidden ₂	Weight Decay	Activation	Best Val Acc.
Coarse	13	0.03	512	256	0	ReLU	0.6600
Coarse	16	0.03	512	256	10^{-4}	ReLU	0.6583
Fine	26	0.06	768	384	10^{-5}	ReLU	0.6711
Fine	16	0.03	768	384	0	ReLU	0.6709
Fine	17	0.03	768	384	10^{-5}	ReLU	0.6706

Figures 1 and 2 summarize the most important trends from the search space, while the full ranking plots for the coarse and fine searches are moved to Appendix A. Several observations are clear. First, larger hidden widths consistently help. Second, ReLU is superior to Tanh in the coarse stage, and both are substantially stronger than Sigmoid. Third, learning rates in the range 0.03 to 0.06 are clearly better than smaller values such as 0.003. Finally, within the explored range, weight decay has a relatively small effect compared with the learning rate and model width.

6 Experimental Results

6.1 Validation and Test Performance

After selecting the best configuration through the two-stage hyperparameter search, the model was trained for a final 30-epoch run and then evaluated on the held-out test set. This final experiment achieved a best validation accuracy of 0.6788 and a final test accuracy of 0.6780, as recorded in `results/curves/full_run_history.json` and `results/errors/full_run_test_summary.json`. These values are higher than the earlier single-stage run and confirm that the search-first strategy improves the final performance of the from-scratch MLP.




Figure 1: Grouped coarse-search statistics. The left panel shows that 0.03 is the strongest learning rate among the coarse candidates, while the right panel shows that larger hidden dimensions consistently improve validation accuracy.




Figure 2: Additional evidence from the hyperparameter study. The left panel shows the coarse-search comparison among ReLU, Tanh, and Sigmoid, where ReLU performs best and Sigmoid is clearly weakest on average. The right panel is the fine-search heatmap over learning rate and first-layer hidden width, showing that wider models combined with learning rates around 0.03 to 0.06 are consistently strong.

The final model used the best fine-search configuration:

- learning rate: 0.06
- hidden dimensions: 768 and 384
- weight decay: 10^{-5}
- activation: ReLU
- training epochs: 30

Table 6: Final experiment summary from the completed two-stage pipeline.

Item	Value
Coarse search trials	54
Fine search trials	27
Final training epochs	30
Best validation accuracy	0.6788
Best validation epoch	30
Test accuracy	0.6780
Best final configuration	LR = 0.06, hidden = 768/384, WD = 10^{-5} , ReLU

Figure 3 shows the corresponding training dynamics for this final run. Training loss decreased from 1.7484 to 0.2422, while validation accuracy improved from 0.4778 to 0.6788. The best validation epoch occurred at the end of training, which suggests that the chosen schedule was still productive and did not yet exhibit a severe performance drop due to overfitting.




Figure 3: Training loss, validation loss, and validation accuracy for the final 30-epoch run after hyperparameter selection. The validation accuracy improves steadily and reaches its best value at the end of training.

The confusion matrix in Figure 4 provides a more detailed class-wise view. Some classes are recognized relatively well:

- Industrial: 0.8347 class accuracy
- Forest: 0.7956 class accuracy
- SeaLake: 0.8333 class accuracy
- Pasture: 0.7700 class accuracy
- Residential: 0.7022 class accuracy
- River: 0.6853 class accuracy

However, other categories remain difficult:

- Highway: 0.4027 class accuracy
- HerbaceousVegetation: 0.5178 class accuracy
- PermanentCrop: 0.5947 class accuracy

Several confusion patterns are especially noticeable. Highway is still frequently confused with River and Residential, although its class accuracy improves relative to the earlier run. PermanentCrop remains strongly entangled with AnnualCrop and HerbaceousVegetation, which reflects the visual similarity of agricultural texture and vegetation distribution. Forest and SeaLake are both strong classes overall, but Forest is still occasionally confused with SeaLake in images dominated by large, dark, low-texture regions.

6.2 First-Layer Weight Visualization

The first-layer weight visualization is included in Appendix B. The weights were reshaped back to the spatial image size for qualitative inspection. Compared with convolutional filters, these patterns are much less interpretable. Most units show diffuse color-sensitive mosaics rather than crisp localized edges or object-part templates.

This behavior is consistent with the architecture. Because the model is a fully connected MLP operating on flattened pixels, it does not preserve local spatial neighborhoods in the same way as a convolutional network. Even so, the weight maps are not completely structureless: several units exhibit weak preferences for greenish vegetation tones, bluish water-like tones, and brown-gray man-made surfaces. This suggests that the first layer has learned coarse global color and texture cues, although the representation is still limited in spatial specificity.




Figure 4: Confusion matrix on the test set for the final 30-epoch run. Strong diagonal entries indicate correct predictions, while off-diagonal mass highlights systematic class confusions.

6.3 Error Analysis

Representative misclassified test samples are included in Appendix B. The selected examples match the confusion-matrix trends and reveal several recurring failure modes.

First, linear structures create ambiguity among Highway, River, and Residential. In multiple examples, a wide road corridor or road-like strip is predicted as River, while dense urban or industrial layouts are sometimes predicted as Highway or Residential. This indicates that the MLP is sensitive to coarse global geometry but does not robustly distinguish material type or surrounding context.

Second, agricultural classes are visually entangled. AnnualCrop, PermanentCrop, HerbaceousVegetation, and Pasture often share repetitive field textures, similar color tones, and large homogeneous regions. When the image is flattened, the model must rely on global appearance rather than local arrangement, which makes these categories particularly difficult to separate.

Third, dark and homogeneous images can also be confusing. For example, SeaLake and Forest can both contain large low-texture regions, and some Forest images are misclassified as Pasture or HerbaceousVegetation when canopy texture is weak.

Overall, the error cases suggest that the model has learned more robust global representations than in the earlier baseline, but it still lacks an explicit mechanism for modeling translation-invariant local structure. This limitation is expected for a dense MLP on remote-sensing imagery and explains why semantically related classes remain hard to distinguish even after careful hyperparameter tuning.

7 Conclusion

This homework implemented a complete three-layer MLP classifier for EuroSAT RGB classification using only NumPy and manual backpropagation. The project covers the full pipeline required by the assignment: data loading, preprocessing, model definition, loss computation, SGD optimization,

learning-rate decay, weight decay, checkpoint saving, staged hyperparameter search, testing, and visualization.

The final experiment used a clear strategy: broad coarse search, focused fine search, and then full training with the best refined configuration. This strategy raised the final validation accuracy to 0.6788 and the test accuracy to 0.6780, which improves on the earlier baseline and provides a more convincing experimental study for the report.

From the completed experiments, three main conclusions are supported by the data. First, a staged hyperparameter search is worthwhile even for a homework-scale MLP: the final system improved over the earlier baseline because search was treated as a real experiment instead of a small appendix. Second, model width and learning rate matter more than small changes in weight decay within the tested range, and the coarse activation comparison shows that Sigmoid is notably less suitable than ReLU or Tanh for this optimization setting. Third, the remaining errors are not random; they mostly reflect visually similar remote-sensing categories that are difficult for a fully connected model operating on flattened pixels.

The main strengths of the project are the correctness of the from-scratch implementation, the modular experiment pipeline, and the more complete hyperparameter study. The main weakness remains representational capacity: without convolution, the model still cannot exploit local spatial structure as effectively as a CNN. A natural extension would be to preserve the manually implemented learning framework while exploring architectures that better model local patterns or spatial inductive bias.

Acknowledgments

This report is based on the implemented homework repository, the provided EuroSAT RGB dataset, the generated experiment artifacts under `results/`, and the assignment handout `hw1.pdf`.

References

- [1] Course handout for Homework 1, provided locally as `hw1.pdf`.
- [2] Project implementation and experiment artifacts under `code/` and `results/`.

Reproducibility Notes

The report text above is tied to the currently available local experiment artifacts. In particular, the cited values come from:

- `results/curves/full_run_history.json`
- `results/errors/full_run_test_summary.json`
- `results/search/full_run_coarse_search_summary.json`
- `results/search/full_run_fine_search_summary.json`
- the corresponding PNG figures under `results/curves/`, `results/errors/`, `results/search/`, and `results/weights/`

If the experiments are rerun and these result files change, the numerical values and discussion in the report should be updated accordingly.

A Full Search Ranking Figures

The main text focuses on grouped statistics and interaction heatmaps because they convey the search trends more compactly. For completeness, the full per-configuration ranking plots for the coarse and fine searches are included below.




Figure 5: Full ranking of validation accuracy for the coarse search configurations. The highest-ranked settings are dominated by ReLU, larger hidden dimensions, and learning rate 0.03, while Sigmoid configurations are concentrated in the lower part of the ranking.




Figure 6: Full ranking of validation accuracy for the fine search configurations. The best-performing region is concentrated around hidden widths 768/384 and learning rates 0.03 to 0.06.

B Additional Qualitative Figure




Figure 7: Visualization of first-layer weights reshaped to the input image size. The patterns are mostly coarse color-texture templates rather than sharp local filters, reflecting the limitations of a flattened MLP representation.




Figure 8: Representative misclassified test samples. Many errors occur between visually similar land-cover categories such as Highway versus River/Residential and AnnualCrop versus PermanentCrop/HerbaceousVegetation.